

VERIFIED RESULTS REPORT

Feedback-lifecycle-audit Held-out TEST-197 / TEST-168

Human reliability, adjudicated lifecycle prevalence, and B/C evaluator comparison

Field	Value
Status	Verified; all reported TEST-set quantities recomputed
Audit date	29 September 2026
Source partition	197 transitions
Analytic cohort	168 frozen claim-bearing packets
Provider calls	0 during verification

Verdict: No false quantitative TEST-set result was detected. The model-specific paired effects are small and point in opposite directions; the supported conclusion is no consistent directional advantage, not statistical equivalence.

1. Cohort and provenance

The report uses TEST-197 for the frozen source partition and TEST-168 for the final claim-bearing analytic cohort. Keeping these names distinct avoids implying that all 197 source transitions received human labels or model decisions.

Stage	Count	Boundary
Source TEST partition	197	Frozen test transitions
Confirmed C	188	Nine objective language exclusions
Stage-A calls	190	50 historical Batch 1 plus 140 Batch 2
Accepted hints	170	Twenty preserved Stage-A failures
Final analytic cohort	168	3 problems, 49 trajectories, 42 participants
Locked state replays	252	S_t and S_t+1 for 126 new packet cases

Source: frozen cohort profile, packet manifest, screening ledger, and integrity summary.

2. Human annotation reliability

Field	Agreement	Raw agreement	Cohen kappa
Original validity	161/168	95.83%	0.2170
Target state at t+1	145/168	86.31%	0.7249
Lifecycle label	145/168	86.31%	0.7249
Instructional priority	144/168	85.71%	0.7117

Pre-adjudication A01/A02 agreement; kappa is unweighted.

All four core fields agreed jointly on 144/168 cases. The 24 unique disagreement cases comprise 23 lifecycle disagreements and one priority-only disagreement.

2.1 Final adjudicated gold

Label	Count	Share	Meaning
KEEP	100	59.52%	Claim remains active
RETIRE	65	38.69%	Supported claim resolved
RETRACT	3	1.79%	Original claim not supported
UNSURE	0	0.00%	Indeterminate

Among 165 hints supported at generation time, 65 were resolved by the next observed revision (39.39%).

Adjudication freeze: All 24 disagreement cases are present in the final JSON and JSONL exports, all required fields are populated, all evidence references are valid, and the blind packet remains unchanged.

3. Frozen primary matrix and predeclared recovery

The primary matrix is complete: 4 frozen model-condition runs and 672 first attempts. Primary files were not overwritten by recovery.

Model	Cond.	Primary 429	Recovery calls	Accepted	Errors
DeepSeek V4.1 Flash	B	90	90	83	7
DeepSeek V4.1 Flash	C	53	53	24	29
Qwen3.8-27B	B	51	51	28	23
Qwen3.8-27B	C	27	27	8	19

Recovery policy: exactly one call per primary HTTP 429; 20 seconds between calls, 90 seconds after another 429; no retry or fallback; non-429 failures excluded.

4. Post-recovery model results

Model	Cond.	Usable	Abstain	Decisions	Correct	Accuracy
DeepSeek	B	154	0	154	146/154	94.81%
DeepSeek	C	135	3	132	127/132	96.21%
Qwen	B	136	0	136	133/136	97.79%
Qwen	C	141	4	137	132/137	96.35%

Decision accuracy is calculated on non-abstaining outputs under the frozen shared evaluator. Condition B is normalized as non-abstaining because it has no selective-decision field.

4.1 Safety-relevant lifecycle errors

Model	Cond.	False KEEP	KEEP recall	Wrongful removal	RETRACT recall
DeepSeek	B	1/68	87/100	6/100	3/3
DeepSeek	C	0/68	73/100	5/100	3/3
Qwen	B	0/68	78/100	3/100	2/3
Qwen	C	1/68	76/100	4/100	2/3

5. Paired common-decision comparison

Model	Paired n	Condition B	Condition C	C - B
DeepSeek V4.1 Flash	124	117/124 (94.35%)	119/124 (95.97%)	+1.61 pp
Qwen3.8-27B	115	113/115 (98.26%)	111/115 (96.52%)	-1.74 pp

Descriptively aggregating the model-specific paired sets gives 239 model-case comparisons. B and C are each correct on 230/239 (96.23%). Both conditions are simultaneously correct on the same case in 228 comparisons.

Each condition produces one false KEEP and eight wrongful removals in the descriptive aggregate. Because model-specific paired subsets differ and cases are clustered, this is not a formal pooled estimator.




Figure 1. Qwen paired common-decision accuracy. The paired subset contains 115 cases.

5.1 Cross-model effect direction




Figure 2. Cross-model paired C - B effect. Positive values favor C; negative values favor B.

Supported conclusion: On the natural held-out cohort, explicit lifecycle structuring does not yield a consistent accuracy advantage over the matched-evidence auditor. DeepSeek moves modestly toward C; Qwen moves modestly toward B.

6. Interpretation and reporting boundary

The paired point estimates are descriptive. No paired confidence interval, equivalence test, or null-hypothesis test is included, so the report does not claim statistical equivalence or proof of no effect.

Rare-label inference remains limited because the natural gold contains only three RETRACT and no UNSURE cases. Controlled Stress-20 results should remain in a separate diagnostic section and must not alter natural prevalence estimates.

Operational coverage and decision quality are distinct. Provider failures remain retained in the ledgers, while accepted explicit abstentions are excluded from decision-accuracy denominators under the frozen evaluator.

7. Reproducibility map

Artifact	Repository path
Human exports	results/heldout_test_197/annotations_frozen_v1/
Blind adjudication packet	results/heldout_test_197/adjudication_24_blind_v1/
Final adjudication	results/heldout_test_197/adjudication_24_final_v1/
Primary and recovery ledgers	results/heldout_test_197/evaluations/
Machine-readable audit	results/heldout_test_197/analysis/verified_test_results.json
Recomputation script	scripts/audit_heldout_test197_results.py

All reported values in this report are generated from the machine-readable audit output. No model predictions were exposed to the adjudicator packet.